

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import os
from scipy import stats
from scipy.stats import norm
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import PolynomialFeatures
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
import sklearn.metrics as metrics
from sklearn import preprocessing
from numpy import array
from sklearn.model_selection import KFold, cross_val_score
from sklearn.linear_model import Lasso
from sklearn.metrics import mean_squared_error
from sklearn.tree import DecisionTreeRegressor

```

```

%matplotlib inline
sns.set()

```

```

df_train = pd.read_csv("train.csv")
df_test = pd.read_csv("test.csv")

```

```

print(df_train.shape)
print("*"*50)
print(df_test.shape)

```

```

(1460, 81)
*****
(1459, 80)

```

```

df_train.head()

```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape
0	1	60	RL	65.0	8450	Pave	NaN	Reg
1	2	20	RL	80.0	9600	Pave	NaN	Reg
2	3	60	RL	68.0	11250	Pave	NaN	IR1
3	4	70	RL	60.0	9550	Pave	NaN	IR1
4	5	60	RL	84.0	14260	Pave	NaN	IR1

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal
MoSold \								
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2								
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
5								
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0
9								
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2								
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0
12								

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

df\_test.head()

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley
LotShape \							
0	1461	20	RH	80.0	11622	Pave	NaN
Reg							
1	1462	20	RL	81.0	14267	Pave	NaN
IR1							
2	1463	60	RL	74.0	13830	Pave	NaN
IR1							
3	1464	60	RL	78.0	9978	Pave	NaN
IR1							
4	1465	120	RL	43.0	5005	Pave	NaN
IR1							

	LandContour	Utilities	...	ScreenPorch	PoolArea	PoolQC	Fence
MiscFeature \							
0	Lvl	AllPub	...	120	0	NaN	MnPrv
NaN							
1	Lvl	AllPub	...	0	0	NaN	NaN
Gar2							
2	Lvl	AllPub	...	0	0	NaN	MnPrv
NaN							
3	Lvl	AllPub	...	0	0	NaN	NaN
NaN							
4	HLS	AllPub	...	144	0	NaN	NaN
NaN							

	MiscVal	MoSold	YrSold	SaleType	SaleCondition
0	0	6	2010	WD	Normal
1	12500	6	2010	WD	Normal
2	0	3	2010	WD	Normal
3	0	6	2010	WD	Normal
4	0	1	2010	WD	Normal

[5 rows x 80 columns]

## EDA

df\_train.describe()

	Id	MSSubClass	LotFrontage	LotArea
OverallQual \				
count	1460.000000	1460.000000	1201.000000	1460.000000
1460.000000				
mean	730.500000	56.897260	70.049958	10516.828082
6.099315				
std	421.610009	42.300571	24.284752	9981.264932
1.382997				
min	1.000000	20.000000	21.000000	1300.000000
1.000000				
25%	365.750000	20.000000	59.000000	7553.500000
5.000000				
50%	730.500000	50.000000	69.000000	9478.500000
6.000000				
75%	1095.250000	70.000000	80.000000	11601.500000
7.000000				
max	1460.000000	190.000000	313.000000	215245.000000
10.000000				

	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea
BsmtFinSF1 ... \				
count	1460.000000	1460.000000	1460.000000	1452.000000
1460.000000 ...				
mean	5.575342	1971.267808	1984.865753	103.685262
443.639726 ...				
std	1.112799	30.202904	20.645407	181.066207
456.098091 ...				
min	1.000000	1872.000000	1950.000000	0.000000
0.000000 ...				
25%	5.000000	1954.000000	1967.000000	0.000000
0.000000 ...				
50%	5.000000	1973.000000	1994.000000	0.000000
383.500000 ...				
75%	6.000000	2000.000000	2004.000000	166.000000
712.250000 ...				
max	9.000000	2010.000000	2010.000000	1600.000000

5644.000000 ...

	WoodDeckSF	OpenPorchSF	EnclosedPorch	3SsnPorch
ScreenPorch \				
count	1460.000000	1460.000000	1460.000000	1460.000000
1460.000000				
mean	94.244521	46.660274	21.954110	3.409589
15.060959				
std	125.338794	66.256028	61.119149	29.317331
55.757415				
min	0.000000	0.000000	0.000000	0.000000
0.000000				
25%	0.000000	0.000000	0.000000	0.000000
0.000000				
50%	0.000000	25.000000	0.000000	0.000000
0.000000				
75%	168.000000	68.000000	0.000000	0.000000
0.000000				
max	857.000000	547.000000	552.000000	508.000000
480.000000				

	PoolArea	MiscVal	MoSold	YrSold
SalePrice				
count	1460.000000	1460.000000	1460.000000	1460.000000
1460.000000				
mean	2.758904	43.489041	6.321918	2007.815753
180921.195890				
std	40.177307	496.123024	2.703626	1.328095
79442.502883				
min	0.000000	0.000000	1.000000	2006.000000
34900.000000				
25%	0.000000	0.000000	5.000000	2007.000000
129975.000000				
50%	0.000000	0.000000	6.000000	2008.000000
163000.000000				
75%	0.000000	0.000000	8.000000	2009.000000
214000.000000				
max	738.000000	15500.000000	12.000000	2010.000000
755000.000000				

[8 rows x 38 columns]

df\_test.describe()

	Id	MSSubClass	LotFrontage	LotArea
OverallQual \				
count	1459.000000	1459.000000	1232.000000	1459.000000
1459.000000				
mean	2190.000000	57.378341	68.580357	9819.161069
6.078821				

std	421.321334	42.746880	22.376841	4955.517327
1.436812				
min	1461.000000	20.000000	21.000000	1470.000000
1.000000				
25%	1825.500000	20.000000	58.000000	7391.000000
5.000000				
50%	2190.000000	50.000000	67.000000	9399.000000
6.000000				
75%	2554.500000	70.000000	80.000000	11517.500000
7.000000				
max	2919.000000	190.000000	200.000000	56600.000000
10.000000				

	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea
BsmtFinSF1	...			
count	1459.000000	1459.000000	1459.000000	1444.000000
1458.000000	...			
mean	5.553804	1971.357779	1983.662783	100.709141
439.203704	...			
std	1.113740	30.390071	21.130467	177.625900
455.268042	...			
min	1.000000	1879.000000	1950.000000	0.000000
0.000000	...			
25%	5.000000	1953.000000	1963.000000	0.000000
0.000000	...			
50%	5.000000	1973.000000	1992.000000	0.000000
350.500000	...			
75%	6.000000	2001.000000	2004.000000	164.000000
753.500000	...			
max	9.000000	2010.000000	2010.000000	1290.000000
4010.000000	...			

	GarageArea	WoodDeckSF	OpenPorchSF	EnclosedPorch
3SsnPorch	...			
count	1458.000000	1459.000000	1459.000000	1459.000000
1459.000000				
mean	472.768861	93.174777	48.313914	24.243317
1.794380				
std	217.048611	127.744882	68.883364	67.227765
20.207842				
min	0.000000	0.000000	0.000000	0.000000
0.000000				
25%	318.000000	0.000000	0.000000	0.000000
0.000000				
50%	480.000000	0.000000	28.000000	0.000000
0.000000				
75%	576.000000	168.000000	72.000000	0.000000
0.000000				
max	1488.000000	1424.000000	742.000000	1012.000000
360.000000				

	ScreenPorch	PoolArea	MiscVal	MoSold
YrSold				
count	1459.000000	1459.000000	1459.000000	1459.000000
mean	17.064428	1.744345	58.167923	6.104181
std	56.609763	30.491646	630.806978	2.722432
min	0.000000	0.000000	0.000000	1.000000
25%	0.000000	0.000000	0.000000	4.000000
50%	0.000000	0.000000	0.000000	6.000000
75%	0.000000	0.000000	0.000000	8.000000
max	576.000000	800.000000	17000.000000	12.000000

[8 rows x 37 columns]

df\_train.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1460 non-null   int64
1   MSSubClass            1460 non-null   int64
2   MSZoning              1460 non-null   object
3   LotFrontage          1201 non-null   float64
4   LotArea              1460 non-null   int64
5   Street               1460 non-null   object
6   Alley                91 non-null     object
7   LotShape             1460 non-null   object
8   LandContour          1460 non-null   object
9   Utilities            1460 non-null   object
10  LotConfig            1460 non-null   object
11  LandSlope            1460 non-null   object
12  Neighborhood          1460 non-null   object
13  Condition1           1460 non-null   object
14  Condition2           1460 non-null   object
15  BldgType             1460 non-null   object
16  HouseStyle           1460 non-null   object
17  OverallQual          1460 non-null   int64
18  OverallCond          1460 non-null   int64
19  YearBuilt            1460 non-null   int64
20  YearRemodAdd         1460 non-null   int64
```

21	RoofStyle	1460	non-null	object
22	RoofMatl	1460	non-null	object
23	Exterior1st	1460	non-null	object
24	Exterior2nd	1460	non-null	object
25	MasVnrType	1452	non-null	object
26	MasVnrArea	1452	non-null	float64
27	ExterQual	1460	non-null	object
28	ExterCond	1460	non-null	object
29	Foundation	1460	non-null	object
30	BsmtQual	1423	non-null	object
31	BsmtCond	1423	non-null	object
32	BsmtExposure	1422	non-null	object
33	BsmtFinType1	1423	non-null	object
34	BsmtFinSF1	1460	non-null	int64
35	BsmtFinType2	1422	non-null	object
36	BsmtFinSF2	1460	non-null	int64
37	BsmtUnfSF	1460	non-null	int64
38	TotalBsmtSF	1460	non-null	int64
39	Heating	1460	non-null	object
40	HeatingQC	1460	non-null	object
41	CentralAir	1460	non-null	object
42	Electrical	1459	non-null	object
43	1stFlrSF	1460	non-null	int64
44	2ndFlrSF	1460	non-null	int64
45	LowQualFinSF	1460	non-null	int64
46	GrLivArea	1460	non-null	int64
47	BsmtFullBath	1460	non-null	int64
48	BsmtHalfBath	1460	non-null	int64
49	FullBath	1460	non-null	int64
50	HalfBath	1460	non-null	int64
51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64
53	KitchenQual	1460	non-null	object
54	TotRmsAbvGrd	1460	non-null	int64
55	Functional	1460	non-null	object
56	Fireplaces	1460	non-null	int64
57	FireplaceQu	770	non-null	object
58	GarageType	1379	non-null	object
59	GarageYrBlt	1379	non-null	float64
60	GarageFinish	1379	non-null	object
61	GarageCars	1460	non-null	int64
62	GarageArea	1460	non-null	int64
63	GarageQual	1379	non-null	object
64	GarageCond	1379	non-null	object
65	PavedDrive	1460	non-null	object
66	WoodDeckSF	1460	non-null	int64
67	OpenPorchSF	1460	non-null	int64
68	EnclosedPorch	1460	non-null	int64
69	3SsnPorch	1460	non-null	int64
70	ScreenPorch	1460	non-null	int64

```

71 PoolArea      1460 non-null    int64
72 PoolQC        7 non-null      object
73 Fence         281 non-null    object
74 MiscFeature   54 non-null      object
75 MiscVal       1460 non-null    int64
76 MoSold        1460 non-null    int64
77 YrSold        1460 non-null    int64
78 SaleType      1460 non-null    object
79 SaleCondition 1460 non-null    object
80 SalePrice     1460 non-null    int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

```
df_test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1459 entries, 0 to 1458
Data columns (total 80 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1459 non-null   int64
1   MSSubClass            1459 non-null   int64
2   MSZoning              1455 non-null   object
3   LotFrontage          1232 non-null   float64
4   LotArea              1459 non-null   int64
5   Street               1459 non-null   object
6   Alley               107 non-null    object
7   LotShape             1459 non-null   object
8   LandContour          1459 non-null   object
9   Utilities            1457 non-null   object
10  LotConfig            1459 non-null   object
11  LandSlope            1459 non-null   object
12  Neighborhood          1459 non-null   object
13  Condition1           1459 non-null   object
14  Condition2           1459 non-null   object
15  BldgType             1459 non-null   object
16  HouseStyle           1459 non-null   object
17  OverallQual          1459 non-null   int64
18  OverallCond          1459 non-null   int64
19  YearBuilt            1459 non-null   int64
20  YearRemodAdd         1459 non-null   int64
21  RoofStyle            1459 non-null   object
22  RoofMatl            1459 non-null   object
23  Exterior1st         1458 non-null   object
24  Exterior2nd         1458 non-null   object
25  MasVnrType          1443 non-null   object
26  MasVnrArea          1444 non-null   float64
27  ExterQual           1459 non-null   object
28  ExterCond           1459 non-null   object
29  Foundation           1459 non-null   object
30  BsmtQual            1415 non-null   object

```



31	BsmtCond	1414	non-null	object
32	BsmtExposure	1415	non-null	object
33	BsmtFinType1	1417	non-null	object
34	BsmtFinSF1	1458	non-null	float64
35	BsmtFinType2	1417	non-null	object
36	BsmtFinSF2	1458	non-null	float64
37	BsmtUnfSF	1458	non-null	float64
38	TotalBsmtSF	1458	non-null	float64
39	Heating	1459	non-null	object
40	HeatingQC	1459	non-null	object
41	CentralAir	1459	non-null	object
42	Electrical	1459	non-null	object
43	1stFlrSF	1459	non-null	int64
44	2ndFlrSF	1459	non-null	int64
45	LowQualFinSF	1459	non-null	int64
46	GrLivArea	1459	non-null	int64
47	BsmtFullBath	1457	non-null	float64
48	BsmtHalfBath	1457	non-null	float64
49	FullBath	1459	non-null	int64
50	HalfBath	1459	non-null	int64
51	BedroomAbvGr	1459	non-null	int64
52	KitchenAbvGr	1459	non-null	int64
53	KitchenQual	1458	non-null	object
54	TotRmsAbvGrd	1459	non-null	int64
55	Functional	1457	non-null	object
56	Fireplaces	1459	non-null	int64
57	FireplaceQu	729	non-null	object
58	GarageType	1383	non-null	object
59	GarageYrBlt	1381	non-null	float64
60	GarageFinish	1381	non-null	object
61	GarageCars	1458	non-null	float64
62	GarageArea	1458	non-null	float64
63	GarageQual	1381	non-null	object
64	GarageCond	1381	non-null	object
65	PavedDrive	1459	non-null	object
66	WoodDeckSF	1459	non-null	int64
67	OpenPorchSF	1459	non-null	int64
68	EnclosedPorch	1459	non-null	int64
69	3SsnPorch	1459	non-null	int64
70	ScreenPorch	1459	non-null	int64
71	PoolArea	1459	non-null	int64
72	PoolQC	3	non-null	object
73	Fence	290	non-null	object
74	MiscFeature	51	non-null	object
75	MiscVal	1459	non-null	int64
76	MoSold	1459	non-null	int64
77	YrSold	1459	non-null	int64
78	SaleType	1458	non-null	object
79	SaleCondition	1459	non-null	object

```
dtypes: float64(11), int64(26), object(43)
memory usage: 912.0+ KB
```


```
df_train['SalePrice'].describe()
```

```
count      1460.000000
mean       180921.195890
std        79442.502883
min        34900.000000
25%       129975.000000
50%       163000.000000
75%       214000.000000
max        755000.000000
Name: SalePrice, dtype: float64
```

```
sns.distplot(df_train['SalePrice']);
print("Skewness: %f" % df_train['SalePrice'].skew())
print("Kurtosis: %f" % df_train['SalePrice'].kurt())
```


```
C:\Users\16095\anaconda3\lib\site-packages\seaborn\
distributions.py:2557: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar
flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Skewness: 1.882876
Kurtosis: 6.536282
```



```
plt.figure(figsize=(20,6))
sns.heatmap(df_train.isnull(),yticklabels=False,cbar=True,cmap='mako')
```

<AxesSubplot:>



```
total_null = df_train.isnull().sum().sort_values(ascending=False)
#First sum and order all null values for each variable
percentage =
(df_train.isnull().sum()/df_train.isnull().count()).sort_values(ascending=False) #Get the percentage
missing_data = pd.concat([total_null, percentage], axis=1,
keys=['Total', 'Percentage'])
missing_data.head(20)
```

	Total	Percentage
PoolQC	1453	0.995205
MiscFeature	1406	0.963014
Alley	1369	0.937671
Fence	1179	0.807534
FireplaceQu	690	0.472603
LotFrontage	259	0.177397
GarageYrBlt	81	0.055479
GarageCond	81	0.055479
GarageType	81	0.055479
GarageFinish	81	0.055479
GarageQual	81	0.055479
BsmtFinType2	38	0.026027
BsmtExposure	38	0.026027
BsmtQual	37	0.025342
BsmtCond	37	0.025342
BsmtFinType1	37	0.025342
MasVnrArea	8	0.005479
MasVnrType	8	0.005479
Electrical	1	0.000685
Id	0	0.000000

## Categorical

```
categ_vars_ls = ['PoolQC', 'MiscFeature', 'Alley', 'Fence',  
                'FireplaceQu', 'GarageType', 'GarageFinish',  
                'GarageQual',  
                'GarageCond', 'BsmtQual', 'BsmtCond', 'BsmtExposure',  
                'BsmtFinType1', 'BsmtFinType2', 'MasVnrType']
```

```
# Clean train set
```

```
for var in categ_vars_ls:  
    df_train[var].fillna('None', inplace=True)
```

```
# Clean test set
```

```
for var in categ_vars_ls:  
    df_test[var].fillna('None', inplace=True)
```

## Numerical

```
num_vars_ls = ['GarageArea', 'GarageCars', 'BsmtFinSF1', 'BsmtFinSF2',  
              'BsmtUnfSF', 'TotalBsmtSF', 'BsmtFullBath',  
              'BsmtHalfBath',  
              'MasVnrArea']
```

```
# Clean train set
```

```
for var in num_vars_ls:  
    df_train[var].fillna(0, inplace=True)
```

```
# Clean test set
```

```
for var in num_vars_ls:  
    df_test[var].fillna(0, inplace=True)
```

```
vars_ls1 = ['Functional', 'MSZoning', 'Electrical', 'KitchenQual',  
            'Exterior1st',  
            'Exterior2nd', 'SaleType', 'Utilities']
```

```
imputer = SimpleImputer(strategy='most_frequent')
```

```
# Clean train set
```

```
df_train[vars_ls1] =  
pd.DataFrame(imputer.fit_transform(df_train[vars_ls1]),  
index=df_train.index)
```

```
# Clean test set
```

```
df_test[vars_ls1] =  
pd.DataFrame(imputer.fit_transform(df_test[vars_ls1]),  
index=df_test.index)
```

```

train_average_house_neighb = df_train.groupby('Neighborhood')
['LotFrontage']
test_average_house_neighb = df_test.groupby('Neighborhood')
['LotFrontage']

# Clean train set
df_train['LotFrontage'].fillna(train_average_house_neighb.transform(lambda x: x.fillna(x.mean()))), inplace=True)

# Clean test set
df_test['LotFrontage'].fillna(test_average_house_neighb.transform(lambda x: x.fillna(x.mean()))), inplace=True)

# Clean train set
df_train['GarageYrBlt'] =
df_train['GarageYrBlt'].fillna(df_train['YearBuilt'])

# Clean test set
df_test['GarageYrBlt'] =
df_test['GarageYrBlt'].fillna(df_test['YearBuilt'])


## NA Check: Verify that we covered all 'NAs' in our data
print(f'Number of NAs in train df: {sum(df_train.isnull().sum())}')
print(f'Number of NAs in test df: {sum(df_test.isnull().sum())}')

Number of NAs in train df: 0
Number of NAs in test df: 0

plt.figure(figsize=(20,6))
sns.heatmap(df_train.isnull(),yticklabels=False,cbar=True,cmap='mako')

<AxesSubplot:>

```



## Investigate potential features & outliers

Below, We can see a few of the highest correlating predictors of SalePrice. Based on these features, it is obvious that usable square footage cumulatively amounts to the highest correlation to SalePrice (GrLivArea, TotalBsmtSF, 1stFlrSF, GarageArea). Other discrete and categorical variables (OverallQual, GarageCars, FullBath, TotRmsAbvGrd) influence the dependent variable as well.

```
corr_mat = df_train.corr().SalePrice.sort_values(ascending=False)
corr_mat.head(10)
```

```
SalePrice      1.000000
OverallQual    0.790982
GrLivArea      0.708624
GarageCars     0.640409
GarageArea     0.623431
TotalBsmtSF    0.613581
1stFlrSF       0.605852
FullBath       0.560664
TotRmsAbvGrd   0.533723
YearBuilt      0.522897
Name: SalePrice, dtype: float64
```

Below we can see the distribution of a few of these variables and assess how outliers may impact the data.

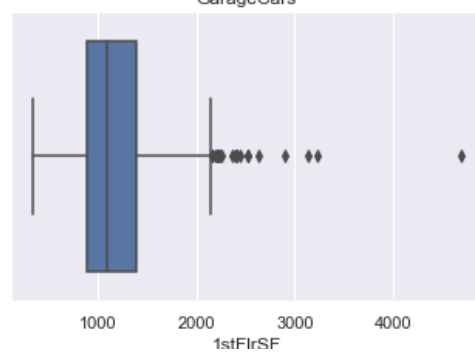





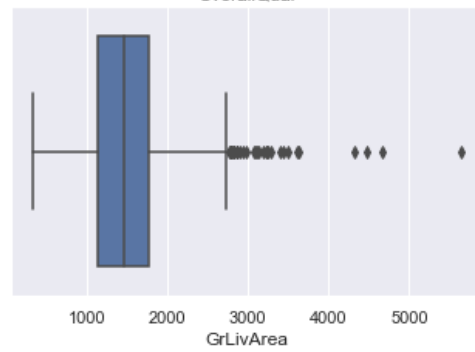



```
cor_features = ['OverallQual', 'GrLivArea', 'TotalBsmtSF',
               'GarageCars', '1stFlrSF', 'YearBuilt' ]

n = len(cor_features)

fig = plt.figure(figsize=(6*2, 4*n))
# add 2 graph for each column variable
gs = fig.add_gridspec(n, 2)
ax = [[fig.add_subplot(gs[i, j]) for j in range(2)] for i in range(n)]

for i in range(n):
    sns.histplot(x=cor_features[i], data=df_train, ax=ax[i][0])
    sns.boxplot(x=cor_features[i], data=df_train, ax=ax[i][1])


plt.show()
```



```

# OverallQual and SalePrice
data = pd.concat([df_train['SalePrice'], df_train['OverallQual']],
axis=1)
f, ax = plt.subplots(figsize=(8, 6))
fig = sns.boxplot(x='OverallQual', y="SalePrice", data=data)
fig.axis(ymin=0, ymax=800000);

```



```

# GrLivArea and SalesPrice
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='GrLivArea', y='SalePrice', data=df_train)
title = plt.title('House Price vs. Above Ground Living Area')

```





The scatter plot above reveals a few outliers where a larger living area is recorded with a low sale price. These outliers can be removed to ensure they do not influence future models.

```
# Clean df_train (GrLiveArea)
outlier = df_train[(df_train.GrLivArea > 4000) & (df_train.SalePrice <
200000)].index
df_train.drop(outlier, axis=0, inplace=True)

# TotalBsmtSF and SalesPrice
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='TotalBsmtSF', y='SalePrice', data=df_train)
title = plt.title('House Price vs. Basement (sqft)')
```



```
# 1stFlrSF and SalePrice
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='1stFlrSF', y='SalePrice', data=df_train)
title = plt.title('House Price vs. First Floor (sqft)')
```



## Feature Creation

Feature creation is likely to be a useful approach to finding more potent predictors in this data set. Based on the list of high correlating variables, it is apparent that features representing usable square feet are strong predictors and can be merged to create a stronger predictive feature. Additionally, the current dataframe seems to categorically discriminate based on above or below ground features. Combining some of high correlation variable, both above and below ground, may yield an overall stronger predictor. Finally, YearBuilt showed up on the bottom of the correlation list with a comparatively low correlation. However, it remains an interesting feature to explore given some obvious and real world implications. Ideally, it would be nice to see in depth how larger renovations might impact the value of older homes. However, the data makes it difficult to define what renovation may have occurred.

Potentially interesting new predictors include:

- Total Square Feet of living Space (Below and Above ground)
- Total Number of Bathrooms (Below and Above Ground)
- Age of House when sold

*# Total Square Feet Column*

```
df_train['TotalSqft'] = df_train['TotalBsmtSF'] + df_train['1stFlrSF']
```

```

+ df_train['2ndFlrSF']

df_test['TotalSqft'] = df_test['TotalBsmtSF'] + df_test['1stFlrSF'] +
df_test['2ndFlrSF']

# Total Bathrooms Column
df_train['TotalBath'] = df_train['FullBath'] +
df_train['BsmtFullBath'] + 0.5*(df_train['HalfBath'] +
df_train['BsmtHalfBath'])

df_test['TotalBath'] = df_test['FullBath'] + df_test['BsmtFullBath'] +
0.5*(df_test['HalfBath'] + df_test['BsmtHalfBath'])

# Age of House
df_train['HouseAge'] = df_train['YrSold'] - df_train['YearBuilt']

df_test['HouseAge'] = df_test['YrSold'] - df_test['YearBuilt']

# Check for new columns
df_train.head()

```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape
0	1	60	RL	65.0	8450	Pave	None	Reg
1	2	20	RL	80.0	9600	Pave	None	Reg
2	3	60	RL	68.0	11250	Pave	None	IR1
3	4	70	RL	60.0	9550	Pave	None	IR1
4	5	60	RL	84.0	14260	Pave	None	IR1

	LandContour	Utilities	...	MiscFeature	MiscVal	MoSold	YrSold
0	Lvl	AllPub	...	None	0	2	2008
1	Lvl	AllPub	...	None	0	5	2007
2	Lvl	AllPub	...	None	0	9	2008
3	Lvl	AllPub	...	None	0	2	2006
4	Lvl	AllPub	...	None	0	12	2008

	SaleCondition	SalePrice	TotalSqft	TotalBath	HouseAge
0	Normal	208500	2566	3.5	5
1	Normal	181500	2524	2.5	31

2	Normal	223500	2706	3.5	7
3	Abnorml	140000	2473	2.0	91
4	Normal	250000	3343	3.5	8

[5 rows x 84 columns]

*# TotalSqft and SalesPrice*

```
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='TotalSqft', y='SalePrice', data=df_train)
title = plt.title('House Price vs. Total Living Space')
```




*# TotalBath and SalesPrice*

```
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='TotalBath', y='SalePrice', data=df_train)
title = plt.title('House Price vs. Total Bathrooms')
```



```
# HouseAge and SalesPrice
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='HouseAge', y='SalePrice', data=df_train)
title = plt.title('House Price vs. Age of House at Sale')
```



```
corr_mat2 = df_train.corr().SalePrice.sort_values(ascending=False)
corr_mat2.head(10)
```

```
SalePrice      1.000000
TotalSqft      0.832877
OverallQual    0.795774
GrLivArea      0.734968
TotalBsmtSF    0.651153
GarageCars     0.641047
TotalBath      0.635896
1stFlrSF       0.631530
GarageArea     0.629217
FullBath       0.562165
Name: SalePrice, dtype: float64
```

## Prep

```
corr_mat = df_train.corr()
corr_mat['SalePrice'][(corr_mat["SalePrice"] > 0.50)]
```

```
OverallQual    0.795774
YearBuilt      0.523608
YearRemodAdd   0.507717
TotalBsmtSF    0.651153
1stFlrSF       0.631530
```

```
GrLivArea      0.734968
FullBath       0.562165
TotRmsAbvGrd   0.537769
GarageYrBlt    0.508719
GarageCars     0.641047
GarageArea     0.629217
SalePrice      1.000000
TotalSqft      0.832877
TotalBath      0.635896
Name: SalePrice, dtype: float64
```

```
important_num_cols = list(corr_mat['SalePrice'][(corr_mat["SalePrice"]
> 0.5)].index)
```

```
important_num_cols.remove('SalePrice')
len(important_num_cols)
print(important_num_cols)
```


```
['OverallQual', 'YearBuilt', 'YearRemodAdd', 'TotalBsmtSF',
'1stFlrSF', 'GrLivArea', 'FullBath', 'TotRmsAbvGrd', 'GarageYrBlt',
'GarageCars', 'GarageArea', 'TotalSqft', 'TotalBath']
```

```
x_num_only = df_train[important_num_cols]
x_num_only.shape
```

```
(1458, 13)
```

```
plt.figure(figsize=(10,8))
sns.heatmap(x_num_only.corr(), center = 0)
plt.title("Correlations Between Columns")
plt.show()
```





```
corr_x = x_num_only.corr()
```

```
for i in range(0, len(corr_x) - 1):
    for j in range(i + 1, len(corr_x)):
        if(corr_x.iloc[i, j] < -0.6 or corr_x.iloc[i, j] > 0.6):
            print(f"corr: {corr_x.iloc[i, j]}, row: {i}, column: {j};
{corr_x.index[i]}, {corr_x.index[j]}")
```

```
corr: 0.6007408233586199, row: 0, column: 9; OverallQual, GarageCars
corr: 0.6773246628152239, row: 0, column: 11; OverallQual, TotalSqft
corr: 0.8448398583188307, row: 1, column: 8; YearBuilt, GarageYrBlt
corr: 0.6036359426335939, row: 2, column: 8; YearRemodAdd, GarageYrBlt
corr: 0.8038296279256135, row: 3, column: 4; TotalBsmtSF, 1stFlrSF
corr: 0.8063997413782633, row: 3, column: 11; TotalBsmtSF, TotalSqft
corr: 0.7819168108101016, row: 4, column: 11; 1stFlrSF, TotalSqft
corr: 0.6383784637415104, row: 5, column: 6; GrLivArea, FullBath
corr: 0.8294981976715332, row: 5, column: 7; GrLivArea, TotRmsAbvGrd
corr: 0.8663861141668201, row: 5, column: 11; GrLivArea, TotalSqft
corr: 0.6932148078878585, row: 6, column: 12; FullBath, TotalBath
corr: 0.6785607174896137, row: 7, column: 11; TotRmsAbvGrd, TotalSqft
corr: 0.6195177230785245, row: 8, column: 9; GarageYrBlt, GarageCars
```

```
corr: 0.6030387164134529, row: 8, column: 10; GarageYrBlt, GarageArea
corr: 0.8873044983919188, row: 9, column: 10; GarageCars, GarageArea
corr: 0.608161005563307, row: 11, column: 12; TotalSqft, TotalBath
```

```
# from the information above, we want to drop '1stFlrSF', 'FullBath',  
'TotRmsAbvGrd', and 'GarageArea'
```

```
num_cols_ls = [i for i in x_num_only.columns if i not in ['1stFlrSF',  
'FullBath', 'TotRmsAbvGrd', 'GarageArea']]
```

```
num_cols_ls
```

```
['OverallQual',  
'YearBuilt',  
'YearRemodAdd',  
'TotalBsmtSF',  
'GrLivArea',  
'GarageYrBlt',  
'GarageCars',  
'TotalSqft',  
'TotalBath']
```

```
# Select the important categorical features to use
```

```
cat_cols = ["MSZoning", "Utilities", "BldgType", "Heating",  
            "KitchenQual", "SaleCondition", "LandSlope"]
```

```
columns_full_ls = num_cols_ls + cat_cols
```

```
columns_full_ls
```

```
['OverallQual',  
'YearBuilt',  
'YearRemodAdd',  
'TotalBsmtSF',  
'GrLivArea',  
'GarageYrBlt',  
'GarageCars',  
'TotalSqft',  
'TotalBath',  
'MSZoning',  
'Utilities',  
'BldgType',  
'Heating',  
'KitchenQual',  
'SaleCondition',  
'LandSlope']
```

```
df_train_final = df_train
```

```
df_test_final = df_test
```

```
df_train_final.head()
```

```
  Id  MSSubClass MSZoning  LotFrontage  LotArea Street Alley LotShape  
\
```

0	1	60	RL	65.0	8450	Pave	None	Reg
1	2	20	RL	80.0	9600	Pave	None	Reg
2	3	60	RL	68.0	11250	Pave	None	IR1
3	4	70	RL	60.0	9550	Pave	None	IR1
4	5	60	RL	84.0	14260	Pave	None	IR1

	LandContour	Utilities	...	MiscFeature	MiscVal	MoSold	YrSold
SaleType \							
0	Lvl	AllPub	...	None	0	2	2008
WD							
1	Lvl	AllPub	...	None	0	5	2007
WD							
2	Lvl	AllPub	...	None	0	9	2008
WD							
3	Lvl	AllPub	...	None	0	2	2006
WD							
4	Lvl	AllPub	...	None	0	12	2008
WD							

	SaleCondition	SalePrice	TotalSqft	TotalBath	HouseAge
0	Normal	208500	2566	3.5	5
1	Normal	181500	2524	2.5	31
2	Normal	223500	2706	3.5	7
3	Abnorml	140000	2473	2.0	91
4	Normal	250000	3343	3.5	8

[5 rows x 84 columns]

```
for col in columns_full_ls:
    train = sorted(df_train_final[col].unique().tolist())
    test = sorted(df_test_final[col].unique().tolist())
    total = set(train + test)
    df_train_final[col] = pd.Categorical(df_train_final[col],
categories=total)
    df_test_final[col] = pd.Categorical(df_test_final[col],
categories=total)

df_train_final = pd.get_dummies(df_train_final, columns=cat_cols)
df_test_final = pd.get_dummies(df_test_final, columns=cat_cols)
```

```
df_train_final.head()
```

	Id	MSSubClass	LotFrontage	LotArea	Street	Alley	LotShape	
LandContour \	0	1	60	65.0	8450	Pave	None	Reg
Lvl	1	2	20	80.0	9600	Pave	None	Reg
Lvl	2	3	60	68.0	11250	Pave	None	IR1
Lvl	3	4	70	60.0	9550	Pave	None	IR1
Lvl	4	5	60	84.0	14260	Pave	None	IR1

	LotConfig	Neighborhood	...	KitchenQual_TA	SaleCondition_Family	\
0	Inside	CollgCr	...	0	0	
1	FR2	Veenker	...	1	0	
2	Inside	CollgCr	...	0	0	
3	Corner	Crawfor	...	0	0	
4	FR2	NoRidge	...	0	0	

	SaleCondition_Abnorml	SaleCondition_Normal	SaleCondition_AdjLand	\
0	0	1	0	
1	0	1	0	
2	0	1	0	
3	1	0	0	
4	0	1	0	

	SaleCondition_Partial	SaleCondition_Alloca	LandSlope_Sev
LandSlope_Mod \	0	0	0
0	0	0	0
1	0	0	0
0	0	0	0
2	0	0	0
0	0	0	0
3	0	0	0
0	0	0	0
4	0	0	0
0			

	LandSlope_Gtl
0	1
1	1
2	1
3	1
4	1

[5 rows x 108 columns]

## Simple linear Regression

*# Feature(s) to look at*

```
f1 = ['TotalSqft']
```

*# Run a Linear Regression using the feature(s)*

```
x1 = df_train_final[f1]
```

```
y = df_train_final['SalePrice']
```

*# Split the data*

```
x_train, x_test, y_train, y_test = train_test_split(x1, y)
```

```
x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

```
((1093, 1), (365, 1), (1093,), (365,))
```

*# Set up model*

```
linreg1 = LinearRegression()
```

```
kf = KFold(n_splits=7, shuffle=True)
```

*# Standardize the data*

```
ss = StandardScaler()
```

```
ss_train = ss.fit_transform(x_train)
```

```
ss_test = ss.transform(x_test)
```

*# cross validate*

```
scores = cross_val_score(linreg1, ss_train, y_train, cv=kf)
```

```
print(scores)
```

```
print(f'Mean Score: {scores.mean()}; SD: {scores.std()}')
```

```
linreg1.fit(ss_train, y_train)
```

```
print(f'TRAIN Score: {linreg1.score(ss_train, y_train)}')
```

```
print(f'TEST Score: {linreg1.score(ss_test, y_test)}')
```

```
pred = linreg1.predict(ss_test)
```

```
b, m = np.polynomial.polynomial.polyfit(y_test, pred, 1)
```

```
[0.62748016 0.77416605 0.62688057 0.64889819 0.69084112 0.73128427  
 0.66161388]
```

```
Mean Score: 0.6801663223202056; SD: 0.05140877822295907
```

```
TRAIN Score: 0.6925391178489992
```

```
TEST Score: 0.6925774318501976
```

The mean score is .68 which is not as good as it could be. The training score is not indicate overfitting.

*# Visualize the model results*

```
sns.scatterplot(x=y_test, y=pred, alpha=0.4)
```

```
sns.regplot(x=y_test, y=pred, truncate=True, scatter_kws={'s': 20,  
'alpha':0.3},
```

```
line_kws={'color':'red', 'linewidth': 2})
```


```

sns.lineplot(x=np.unique(y_test), y=np.unique(np.poly1d(b + m *
np.unique(y_test))), linewidth=0.5, color='r')

plt.xlabel("Actual Sale Prices")
plt.ylabel("Predicted Sale Prices")
plt.title("Actual Prices vs Predicted prices [Test Set]")

plt.show()

```



## Multiple Linear Regression

*# Feature(s) to look at*

```
f2 = num_cols_ls
```

*# Run a Linear Regression using the feature(s)*

```
x2 = df_train_final[f2]
```

```
y = df_train_final['SalePrice']
```

*# Split the data*

```
x2_train, x2_test, y2_train, y2_test = train_test_split(x2, y)
```

```
x2_train.shape, x2_test.shape, y2_train.shape, y2_test.shape
```

```
((1093, 9), (365, 9), (1093,), (365,))
```

*# Multi Linear Regression*

```
linreg2 = LinearRegression()
```

```
kf = KFold(n_splits=7, shuffle=True)
```

```
# Standardize the numeric columns
```

```
ss = StandardScaler()
```

```
ss_train2 = ss.fit_transform(x2_train)
```

```
ss_test2 = ss.transform(x2_test)
```

```
scores = cross_val_score(linreg2, ss_train2, y2_train, cv=kf)
```

```
print(scores)
```

```
print(f'Mean Score: {scores.mean()}; SD: {scores.std()}')
```

```
linreg2.fit(ss_train2, y2_train)
```

```
print(f'TRAIN Score: {linreg2.score(ss_train2, y2_train)}')
```

```
print(f'TEST Score: {linreg2.score(ss_test2, y2_test)}')
```

```
pred2 = linreg2.predict(ss_test2)
```

```
b, m = np.polynomial.polynomial.polyfit(y2_test, pred2, 1)
```

```
[0.84800476 0.7755963  0.81937807 0.81932924 0.82138407 0.77017009  
 0.80274092]
```

```
Mean Score: 0.8080862058622865; SD: 0.025483983264919123
```

```
TRAIN Score: 0.8159443440981854
```

```
TEST Score: 0.8186095109797409
```

Based on the scores above, the Multiple Regression model seems to have performed better than the Simple Regression model. It appears to be fitting to the data better than the previous model, but does not indicate overfitting.

```
sns.scatterplot(x=y2_test, y=pred2, alpha=0.4)
```

```
sns.regplot(x=y2_test, y=pred2, truncate=True, scatter_kws={'s': 20,  
'alpha':0.3},
```

```
          line_kws={'color':'red', 'linewidth': 2})
```

```
sns.lineplot(x=np.unique(y2_test), y=np.unique(np.poly1d(b + m *  
np.unique(y2_test))), linewidth=0.5, color='r')
```

```
plt.xlabel("Actual Sale Prices")
```

```
plt.ylabel("Predicted Sale Prices")
```

```
plt.title("Actual Prices vs Predicted prices [Test Set]")
```

```
plt.show()
```



## Conclusion

I am not entirely certain as to why the multiple regression performed better than simple regression in this case. Perhaps the entirety of the numerical list allowed the model to identify more correlative patterns which in turn allowed for a better fit. I expected that the TotalSqft variable would have provided a better fit considering it encompassed several highly correlative variables in one feature. Nevertheless, both models provide a good enough fit to reflect the inputted data.