

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import os
from scipy import stats
from scipy.stats import norm
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import Ridge
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, ElasticNet,
ElasticNetCV
from sklearn.linear_model import LogisticRegression
import sklearn.metrics as metrics
from sklearn import preprocessing
from numpy import array
from sklearn.model_selection import KFold, cross_val_score
from sklearn.linear_model import Lasso, LassoCV
from sklearn.metrics import mean_squared_error
from sklearn.tree import DecisionTreeRegressor

```

```

%matplotlib inline
sns.set()

```

```

df_train = pd.read_csv("train.csv")
df_test = pd.read_csv("test.csv")

```

```

print(df_train.shape)
print("*"*50)
print(df_test.shape)

```

```

(1460, 81)

```

```

*****

```

```

(1459, 80)

```

```

df_train.head()

```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape
0	1	60	RL	65.0	8450	Pave	NaN	Reg
1	2	20	RL	80.0	9600	Pave	NaN	Reg
2	3	60	RL	68.0	11250	Pave	NaN	IR1
3	4	70	RL	60.0	9550	Pave	NaN	IR1

4	5	60	RL	84.0	14260	Pave	NaN	IR1
---	---	----	----	------	-------	------	-----	-----

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal
MoSold \								
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2								
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
5								
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0
9								
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2								
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0
12								

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

df_test.head()

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley
LotShape \							
0	1461	20	RH	80.0	11622	Pave	NaN
Reg							
1	1462	20	RL	81.0	14267	Pave	NaN
IR1							
2	1463	60	RL	74.0	13830	Pave	NaN
IR1							
3	1464	60	RL	78.0	9978	Pave	NaN
IR1							
4	1465	120	RL	43.0	5005	Pave	NaN
IR1							

	LandContour	Utilities	...	ScreenPorch	PoolArea	PoolQC	Fence
MiscFeature \							
0	Lvl	AllPub	...	120	0	NaN	MnPrv
NaN							
1	Lvl	AllPub	...	0	0	NaN	NaN
Gar2							
2	Lvl	AllPub	...	0	0	NaN	MnPrv
NaN							
3	Lvl	AllPub	...	0	0	NaN	NaN
NaN							

4	HLS	AllPub	...	144	0	NaN	NaN
NaN							

	MiscVal	MoSold	YrSold	SaleType	SaleCondition
0	0	6	2010	WD	Normal
1	12500	6	2010	WD	Normal
2	0	3	2010	WD	Normal
3	0	6	2010	WD	Normal
4	0	1	2010	WD	Normal

[5 rows x 80 columns]

EDA

df_train.describe()

	Id	MSSubClass	LotFrontage	LotArea
OverallQual \				
count	1460.000000	1460.000000	1201.000000	1460.000000
1460.000000				
mean	730.500000	56.897260	70.049958	10516.828082
6.099315				
std	421.610009	42.300571	24.284752	9981.264932
1.382997				
min	1.000000	20.000000	21.000000	1300.000000
1.000000				
25%	365.750000	20.000000	59.000000	7553.500000
5.000000				
50%	730.500000	50.000000	69.000000	9478.500000
6.000000				
75%	1095.250000	70.000000	80.000000	11601.500000
7.000000				
max	1460.000000	190.000000	313.000000	215245.000000
10.000000				

	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea
BsmtFinSF1 ... \				
count	1460.000000	1460.000000	1460.000000	1452.000000
1460.000000 ...				
mean	5.575342	1971.267808	1984.865753	103.685262
443.639726 ...				
std	1.112799	30.202904	20.645407	181.066207
456.098091 ...				
min	1.000000	1872.000000	1950.000000	0.000000
0.000000 ...				
25%	5.000000	1954.000000	1967.000000	0.000000
0.000000 ...				
50%	5.000000	1973.000000	1994.000000	0.000000
383.500000 ...				
75%	6.000000	2000.000000	2004.000000	166.000000

```

712.250000 ...
max      9.000000 2010.000000 2010.000000 1600.000000
5644.000000 ...

```

```

      WoodDeckSF  OpenPorchSF  EnclosedPorch  3SsnPorch
ScreenPorch \
count 1460.000000 1460.000000 1460.000000 1460.000000
1460.000000
mean   94.244521  46.660274   21.954110   3.409589
15.060959
std   125.338794  66.256028   61.119149   29.317331
55.757415
min     0.000000   0.000000   0.000000   0.000000
0.000000
25%     0.000000   0.000000   0.000000   0.000000
0.000000
50%     0.000000  25.000000   0.000000   0.000000
0.000000
75%    168.000000  68.000000   0.000000   0.000000
0.000000
max    857.000000 547.000000 552.000000 508.000000
480.000000

```

```

      PoolArea      MiscVal      MoSold      YrSold
SalePrice
count 1460.000000 1460.000000 1460.000000 1460.000000
1460.000000
mean   2.758904  43.489041   6.321918 2007.815753
180921.195890
std   40.177307 496.123024   2.703626   1.328095
79442.502883
min     0.000000   0.000000   1.000000 2006.000000
34900.000000
25%     0.000000   0.000000   5.000000 2007.000000
129975.000000
50%     0.000000   0.000000   6.000000 2008.000000
163000.000000
75%     0.000000   0.000000   8.000000 2009.000000
214000.000000
max    738.000000 15500.000000 12.000000 2010.000000
755000.000000

```

[8 rows x 38 columns]

```
df_test.describe()
```

```

      Id  MSSubClass  LotFrontage      LotArea
OverallQual \
count 1459.000000 1459.000000 1232.000000 1459.000000
1459.000000

```

mean	2190.000000	57.378341	68.580357	9819.161069
6.078821				
std	421.321334	42.746880	22.376841	4955.517327
1.436812				
min	1461.000000	20.000000	21.000000	1470.000000
1.000000				
25%	1825.500000	20.000000	58.000000	7391.000000
5.000000				
50%	2190.000000	50.000000	67.000000	9399.000000
6.000000				
75%	2554.500000	70.000000	80.000000	11517.500000
7.000000				
max	2919.000000	190.000000	200.000000	56600.000000
10.000000				

	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea
BsmtFinSF1	...			
count	1459.000000	1459.000000	1459.000000	1444.000000
1458.000000	...			
mean	5.553804	1971.357779	1983.662783	100.709141
439.203704	...			
std	1.113740	30.390071	21.130467	177.625900
455.268042	...			
min	1.000000	1879.000000	1950.000000	0.000000
0.000000	...			
25%	5.000000	1953.000000	1963.000000	0.000000
0.000000	...			
50%	5.000000	1973.000000	1992.000000	0.000000
350.500000	...			
75%	6.000000	2001.000000	2004.000000	164.000000
753.500000	...			
max	9.000000	2010.000000	2010.000000	1290.000000
4010.000000	...			

	GarageArea	WoodDeckSF	OpenPorchSF	EnclosedPorch
3SsnPorch	...			
count	1458.000000	1459.000000	1459.000000	1459.000000
1459.000000				
mean	472.768861	93.174777	48.313914	24.243317
1.794380				
std	217.048611	127.744882	68.883364	67.227765
20.207842				
min	0.000000	0.000000	0.000000	0.000000
0.000000				
25%	318.000000	0.000000	0.000000	0.000000
0.000000				
50%	480.000000	0.000000	28.000000	0.000000
0.000000				
75%	576.000000	168.000000	72.000000	0.000000
0.000000				

```
max      1488.000000  1424.000000   742.000000   1012.000000
360.000000
```

```
      ScreenPorch      PoolArea      MiscVal      MoSold
YrSold
count  1459.000000  1459.000000  1459.000000  1459.000000
1459.000000
mean     17.064428     1.744345    58.167923    6.104181
2007.769705
std      56.609763    30.491646   630.806978    2.722432
1.301740
min       0.000000     0.000000     0.000000    1.000000
2006.000000
25%       0.000000     0.000000     0.000000    4.000000
2007.000000
50%       0.000000     0.000000     0.000000    6.000000
2008.000000
75%       0.000000     0.000000     0.000000    8.000000
2009.000000
max      576.000000   800.000000  17000.000000   12.000000
2010.000000
```

```
[8 rows x 37 columns]
```

```
df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1460 entries, 0 to 1459
```

```
Data columns (total 81 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	1460 non-null	int64
1	MSSubClass	1460 non-null	int64
2	MSZoning	1460 non-null	object
3	LotFrontage	1201 non-null	float64
4	LotArea	1460 non-null	int64
5	Street	1460 non-null	object
6	Alley	91 non-null	object
7	LotShape	1460 non-null	object
8	LandContour	1460 non-null	object
9	Utilities	1460 non-null	object
10	LotConfig	1460 non-null	object
11	LandSlope	1460 non-null	object
12	Neighborhood	1460 non-null	object
13	Condition1	1460 non-null	object
14	Condition2	1460 non-null	object
15	BldgType	1460 non-null	object
16	HouseStyle	1460 non-null	object
17	OverallQual	1460 non-null	int64
18	OverallCond	1460 non-null	int64

19	YearBuilt	1460	non-null	int64
20	YearRemodAdd	1460	non-null	int64
21	RoofStyle	1460	non-null	object
22	RoofMatl	1460	non-null	object
23	Exterior1st	1460	non-null	object
24	Exterior2nd	1460	non-null	object
25	MasVnrType	1452	non-null	object
26	MasVnrArea	1452	non-null	float64
27	ExterQual	1460	non-null	object
28	ExterCond	1460	non-null	object
29	Foundation	1460	non-null	object
30	BsmtQual	1423	non-null	object
31	BsmtCond	1423	non-null	object
32	BsmtExposure	1422	non-null	object
33	BsmtFinType1	1423	non-null	object
34	BsmtFinSF1	1460	non-null	int64
35	BsmtFinType2	1422	non-null	object
36	BsmtFinSF2	1460	non-null	int64
37	BsmtUnfSF	1460	non-null	int64
38	TotalBsmtSF	1460	non-null	int64
39	Heating	1460	non-null	object
40	HeatingQC	1460	non-null	object
41	CentralAir	1460	non-null	object
42	Electrical	1459	non-null	object
43	1stFlrSF	1460	non-null	int64
44	2ndFlrSF	1460	non-null	int64
45	LowQualFinSF	1460	non-null	int64
46	GrLivArea	1460	non-null	int64
47	BsmtFullBath	1460	non-null	int64
48	BsmtHalfBath	1460	non-null	int64
49	FullBath	1460	non-null	int64
50	HalfBath	1460	non-null	int64
51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64
53	KitchenQual	1460	non-null	object
54	TotRmsAbvGrd	1460	non-null	int64
55	Functional	1460	non-null	object
56	Fireplaces	1460	non-null	int64
57	FireplaceQu	770	non-null	object
58	GarageType	1379	non-null	object
59	GarageYrBlt	1379	non-null	float64
60	GarageFinish	1379	non-null	object
61	GarageCars	1460	non-null	int64
62	GarageArea	1460	non-null	int64
63	GarageQual	1379	non-null	object
64	GarageCond	1379	non-null	object
65	PavedDrive	1460	non-null	object
66	WoodDeckSF	1460	non-null	int64
67	OpenPorchSF	1460	non-null	int64
68	EnclosedPorch	1460	non-null	int64

```

69  3SsnPorch      1460 non-null   int64
70  ScreenPorch    1460 non-null   int64
71  PoolArea       1460 non-null   int64
72  PoolQC         7 non-null      object
73  Fence          281 non-null   object
74  MiscFeature    54 non-null      object
75  MiscVal        1460 non-null   int64
76  MoSold         1460 non-null   int64
77  YrSold         1460 non-null   int64
78  SaleType       1460 non-null   object
79  SaleCondition  1460 non-null   object
80  SalePrice      1460 non-null   int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

```
df_test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1459 entries, 0 to 1458
Data columns (total 80 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1459 non-null   int64
1   MSSubClass            1459 non-null   int64
2   MSZoning              1455 non-null   object
3   LotFrontage          1232 non-null   float64
4   LotArea               1459 non-null   int64
5   Street               1459 non-null   object
6   Alley                107 non-null    object
7   LotShape             1459 non-null   object
8   LandContour          1459 non-null   object
9   Utilities            1457 non-null   object
10  LotConfig            1459 non-null   object
11  LandSlope            1459 non-null   object
12  Neighborhood          1459 non-null   object
13  Condition1           1459 non-null   object
14  Condition2           1459 non-null   object
15  BldgType             1459 non-null   object
16  HouseStyle           1459 non-null   object
17  OverallQual          1459 non-null   int64
18  OverallCond          1459 non-null   int64
19  YearBuilt            1459 non-null   int64
20  YearRemodAdd         1459 non-null   int64
21  RoofStyle            1459 non-null   object
22  RoofMatl            1459 non-null   object
23  Exterior1st          1458 non-null   object
24  Exterior2nd          1458 non-null   object
25  MasVnrType           1443 non-null   object
26  MasVnrArea           1444 non-null   float64
27  ExterQual            1459 non-null   object
28  ExterCond            1459 non-null   object

```


29	Foundation	1459	non-null	object
30	BsmtQual	1415	non-null	object
31	BsmtCond	1414	non-null	object
32	BsmtExposure	1415	non-null	object
33	BsmtFinType1	1417	non-null	object
34	BsmtFinSF1	1458	non-null	float64
35	BsmtFinType2	1417	non-null	object
36	BsmtFinSF2	1458	non-null	float64
37	BsmtUnfSF	1458	non-null	float64
38	TotalBsmtSF	1458	non-null	float64
39	Heating	1459	non-null	object
40	HeatingQC	1459	non-null	object
41	CentralAir	1459	non-null	object
42	Electrical	1459	non-null	object
43	1stFlrSF	1459	non-null	int64
44	2ndFlrSF	1459	non-null	int64
45	LowQualFinSF	1459	non-null	int64
46	GrLivArea	1459	non-null	int64
47	BsmtFullBath	1457	non-null	float64
48	BsmtHalfBath	1457	non-null	float64
49	FullBath	1459	non-null	int64
50	HalfBath	1459	non-null	int64
51	BedroomAbvGr	1459	non-null	int64
52	KitchenAbvGr	1459	non-null	int64
53	KitchenQual	1458	non-null	object
54	TotRmsAbvGrd	1459	non-null	int64
55	Functional	1457	non-null	object
56	Fireplaces	1459	non-null	int64
57	FireplaceQu	729	non-null	object
58	GarageType	1383	non-null	object
59	GarageYrBlt	1381	non-null	float64
60	GarageFinish	1381	non-null	object
61	GarageCars	1458	non-null	float64
62	GarageArea	1458	non-null	float64
63	GarageQual	1381	non-null	object
64	GarageCond	1381	non-null	object
65	PavedDrive	1459	non-null	object
66	WoodDeckSF	1459	non-null	int64
67	OpenPorchSF	1459	non-null	int64
68	EnclosedPorch	1459	non-null	int64
69	3SsnPorch	1459	non-null	int64
70	ScreenPorch	1459	non-null	int64
71	PoolArea	1459	non-null	int64
72	PoolQC	3	non-null	object
73	Fence	290	non-null	object
74	MiscFeature	51	non-null	object
75	MiscVal	1459	non-null	int64
76	MoSold	1459	non-null	int64
77	YrSold	1459	non-null	int64
78	SaleType	1458	non-null	object

```
79 SaleCondition 1459 non-null object
dtypes: float64(11), int64(26), object(43)
memory usage: 912.0+ KB
```


```
df_train['SalePrice'].describe()
```

```
count      1460.000000
mean       180921.195890
std        79442.502883
min        34900.000000
25%        129975.000000
50%        163000.000000
75%        214000.000000
max        755000.000000
Name: SalePrice, dtype: float64
```

```
sns.distplot(df_train['SalePrice']);
print("Skewness: %f" % df_train['SalePrice'].skew())
print("Kurtosis: %f" % df_train['SalePrice'].kurt())
```


```
C:\Users\16095\anaconda3\lib\site-packages\seaborn\
distributions.py:2557: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar
flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Skewness: 1.882876
Kurtosis: 6.536282
```



```
plt.figure(figsize=(20,6))
sns.heatmap(df_train.isnull(),yticklabels=False,cbar=True,cmap='mako')
```

<AxesSubplot:>



```
total_null = df_train.isnull().sum().sort_values(ascending=False)
#First sum and order all null values for each variable
percentage =
(df_train.isnull().sum()/df_train.isnull().count()).sort_values(ascending=False) #Get the percentage
missing_data = pd.concat([total_null, percentage], axis=1,
keys=['Total', 'Percentage'])
missing_data.head(20)
```

	Total	Percentage
PoolQC	1453	0.995205
MiscFeature	1406	0.963014
Alley	1369	0.937671
Fence	1179	0.807534
FireplaceQu	690	0.472603
LotFrontage	259	0.177397
GarageYrBlt	81	0.055479
GarageCond	81	0.055479
GarageType	81	0.055479
GarageFinish	81	0.055479
GarageQual	81	0.055479
BsmtFinType2	38	0.026027
BsmtExposure	38	0.026027
BsmtQual	37	0.025342
BsmtCond	37	0.025342
BsmtFinType1	37	0.025342
MasVnrArea	8	0.005479
MasVnrType	8	0.005479
Electrical	1	0.000685
Id	0	0.000000

Categorical

```
categ_vars_ls = ['PoolQC', 'MiscFeature', 'Alley', 'Fence',  
                'FireplaceQu', 'GarageType', 'GarageFinish',  
                'GarageQual',  
                'GarageCond', 'BsmtQual', 'BsmtCond', 'BsmtExposure',  
                'BsmtFinType1', 'BsmtFinType2', 'MasVnrType']
```

```
# Clean train set
```

```
for var in categ_vars_ls:  
    df_train[var].fillna('None', inplace=True)
```

```
# Clean test set
```

```
for var in categ_vars_ls:  
    df_test[var].fillna('None', inplace=True)
```

Numerical

```
num_vars_ls = ['GarageArea', 'GarageCars', 'BsmtFinSF1', 'BsmtFinSF2',  
              'BsmtUnfSF', 'TotalBsmtSF', 'BsmtFullBath',  
              'BsmtHalfBath',  
              'MasVnrArea']
```

```
# Clean train set
```

```
for var in num_vars_ls:  
    df_train[var].fillna(0, inplace=True)
```

```
# Clean test set
```

```
for var in num_vars_ls:  
    df_test[var].fillna(0, inplace=True)
```

```
vars_ls1 = ['Functional', 'MSZoning', 'Electrical', 'KitchenQual',  
            'Exterior1st',  
            'Exterior2nd', 'SaleType', 'Utilities']
```

```
imputer = SimpleImputer(strategy='most_frequent')
```

```
# Clean train set
```

```
df_train[vars_ls1] =  
pd.DataFrame(imputer.fit_transform(df_train[vars_ls1]),  
index=df_train.index)
```

```
# Clean test set
```

```
df_test[vars_ls1] =  
pd.DataFrame(imputer.fit_transform(df_test[vars_ls1]),  
index=df_test.index)
```

```

train_average_house_neighb = df_train.groupby('Neighborhood')
['LotFrontage']
test_average_house_neighb = df_test.groupby('Neighborhood')
['LotFrontage']

# Clean train set
df_train['LotFrontage'].fillna(train_average_house_neighb.transform(lambda x: x.fillna(x.mean()))), inplace=True)

# Clean test set
df_test['LotFrontage'].fillna(test_average_house_neighb.transform(lambda x: x.fillna(x.mean()))), inplace=True)

# Clean train set
df_train['GarageYrBlt'] =
df_train['GarageYrBlt'].fillna(df_train['YearBuilt'])

# Clean test set
df_test['GarageYrBlt'] =
df_test['GarageYrBlt'].fillna(df_test['YearBuilt'])


## NA Check: Verify that we covered all 'NAs' in our data
print(f'Number of NAs in train df: {sum(df_train.isnull().sum())}')
print(f'Number of NAs in test df: {sum(df_test.isnull().sum())}')

Number of NAs in train df: 0
Number of NAs in test df: 0

plt.figure(figsize=(20,6))
sns.heatmap(df_train.isnull(),yticklabels=False,cbar=True,cmap='mako')

<AxesSubplot:>

```



Investigate potential features & outliers

Below, We can see a few of the highest correlating predictors of SalePrice. Based on these features, it is obvious that usable square footage cumulatively amounts to the highest correlation to SalePrice (GrLivArea, TotalBsmtSF, 1stFlrSF, GarageArea). Other discrete and categorical variables (OverallQual, GarageCars, FullBath, TotRmsAbvGrd) influence the dependent variable as well.

```
corr_mat = df_train.corr().SalePrice.sort_values(ascending=False)
corr_mat.head(10)
```

```
SalePrice      1.000000
OverallQual    0.790982
GrLivArea      0.708624
GarageCars     0.640409
GarageArea     0.623431
TotalBsmtSF    0.613581
1stFlrSF       0.605852
FullBath       0.560664
TotRmsAbvGrd   0.533723
YearBuilt      0.522897
Name: SalePrice, dtype: float64
```

Below we can see the distribution of a few of these variables and assess how outliers may impact the data.











```
cor_features = ['OverallQual', 'GrLivArea', 'TotalBsmtSF',
                'GarageCars', '1stFlrSF', 'YearBuilt' ]

n = len(cor_features)

fig = plt.figure(figsize=(6*2, 4*n))
# add 2 graph for each column variable
gs = fig.add_gridspec(n, 2)
ax = [[fig.add_subplot(gs[i, j]) for j in range(2)] for i in range(n)]

for i in range(n):
    sns.histplot(x=cor_features[i], data=df_train, ax=ax[i][0])
    sns.boxplot(x=cor_features[i], data=df_train, ax=ax[i][1])


plt.show()
```



```

# OverallQual and SalePrice
data = pd.concat([df_train['SalePrice'], df_train['OverallQual']],
axis=1)
f, ax = plt.subplots(figsize=(8, 6))
fig = sns.boxplot(x='OverallQual', y="SalePrice", data=data)
fig.axis(ymin=0, ymax=800000);

```



```

# GrLivArea and SalesPrice
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='GrLivArea', y='SalePrice', data=df_train)
title = plt.title('House Price vs. Above Ground Living Area')

```




The scatter plot above reveals a few outliers where a larger living area is recorded with a low sale price. These outliers can be removed to ensure they do not influence future models.

```
# Clean df_train (GrLiveArea)
outlier = df_train[(df_train.GrLivArea > 4000) & (df_train.SalePrice < 200000)].index
df_train.drop(outlier, axis=0, inplace=True)

# TotalBsmtSF and SalesPrice
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='TotalBsmtSF', y='SalePrice', data=df_train)
title = plt.title('House Price vs. Basement (sqft)')
```



```
# 1stFlrSF and SalePrice
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='1stFlrSF', y='SalePrice', data=df_train)
title = plt.title('House Price vs. First Floor (sqft)')
```



Feature Creation

Feature creation is likely to be a useful approach to finding more potent predictors in this data set. Based on the list of high correlating variables, it is apparent that features representing usable square feet are strong predictors and can be merged to create a stronger predictive feature. Additionally, the current dataframe seems to categorically discriminate based on above or below ground features. Combining some of high correlation variable, both above and below ground, may yield an overall stronger predictor. Finally, YearBuilt showed up on the bottom of the correlation list with a comparatively low correlation. However, it remains an interesting feature to explore given some obvious and real world implications. Ideally, it would be nice to see in depth how larger renovations might impact the value of older homes. However, the data makes it difficult to define what renovation may have occurred.

Potentially interesting new predictors include:

- Total Square Feet of living Space (Below and Above ground)
- Total Number of Bathrooms (Below and Above Ground)
- Age of House when sold

Total Square Feet Column

```
df_train['TotalSqft'] = df_train['TotalBsmtSF'] + df_train['1stFlrSF']
```

```

+ df_train['2ndFlrSF']

df_test['TotalSqft'] = df_test['TotalBsmtSF'] + df_test['1stFlrSF'] +
df_test['2ndFlrSF']

# Total Bathrooms Column
df_train['TotalBath'] = df_train['FullBath'] +
df_train['BsmtFullBath'] + 0.5*(df_train['HalfBath'] +
df_train['BsmtHalfBath'])

df_test['TotalBath'] = df_test['FullBath'] + df_test['BsmtFullBath'] +
0.5*(df_test['HalfBath'] + df_test['BsmtHalfBath'])

# Age of House
df_train['HouseAge'] = df_train['YrSold'] - df_train['YearBuilt']

df_test['HouseAge'] = df_test['YrSold'] - df_test['YearBuilt']

# Check for new columns
df_train.head()

```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape
0	1	60	RL	65.0	8450	Pave	None	Reg
1	2	20	RL	80.0	9600	Pave	None	Reg
2	3	60	RL	68.0	11250	Pave	None	IR1
3	4	70	RL	60.0	9550	Pave	None	IR1
4	5	60	RL	84.0	14260	Pave	None	IR1

	LandContour	Utilities	...	MiscFeature	MiscVal	MoSold	YrSold
0	Lvl	AllPub	...	None	0	2	2008
1	Lvl	AllPub	...	None	0	5	2007
2	Lvl	AllPub	...	None	0	9	2008
3	Lvl	AllPub	...	None	0	2	2006
4	Lvl	AllPub	...	None	0	12	2008

	SaleCondition	SalePrice	TotalSqft	TotalBath	HouseAge
0	Normal	208500	2566	3.5	5
1	Normal	181500	2524	2.5	31

2	Normal	223500	2706	3.5	7
3	Abnorml	140000	2473	2.0	91
4	Normal	250000	3343	3.5	8

[5 rows x 84 columns]

TotalSqft and SalesPrice

```
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='TotalSqft', y='SalePrice', data=df_train)
title = plt.title('House Price vs. Total Living Space')
```




TotalBath and SalesPrice

```
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='TotalBath', y='SalePrice', data=df_train)
title = plt.title('House Price vs. Total Bathrooms')
```



```
# HouseAge and SalesPrice
sns.set_style('darkgrid')
plt.figure(figsize=(8, 6))
sns.scatterplot(x='HouseAge', y='SalePrice', data=df_train)
title = plt.title('House Price vs. Age of House at Sale')
```



```
corr_mat2 = df_train.corr().SalePrice.sort_values(ascending=False)
corr_mat2.head(10)
```

```
SalePrice      1.000000
TotalSqft      0.832877
OverallQual    0.795774
GrLivArea      0.734968
TotalBsmtSF    0.651153
GarageCars     0.641047
TotalBath      0.635896
1stFlrSF       0.631530
GarageArea     0.629217
FullBath       0.562165
Name: SalePrice, dtype: float64
```

Skewness Check (train)

```
numerics = ['int16', 'int32', 'int64', 'float16', 'float32',
'float64']
```

```
dataset_numeric = df_train.select_dtypes(include=numerics)
```

```
dataset_numeric.shape
```

```
(1458, 41)
```

```
dataset_numeric.skew()
```

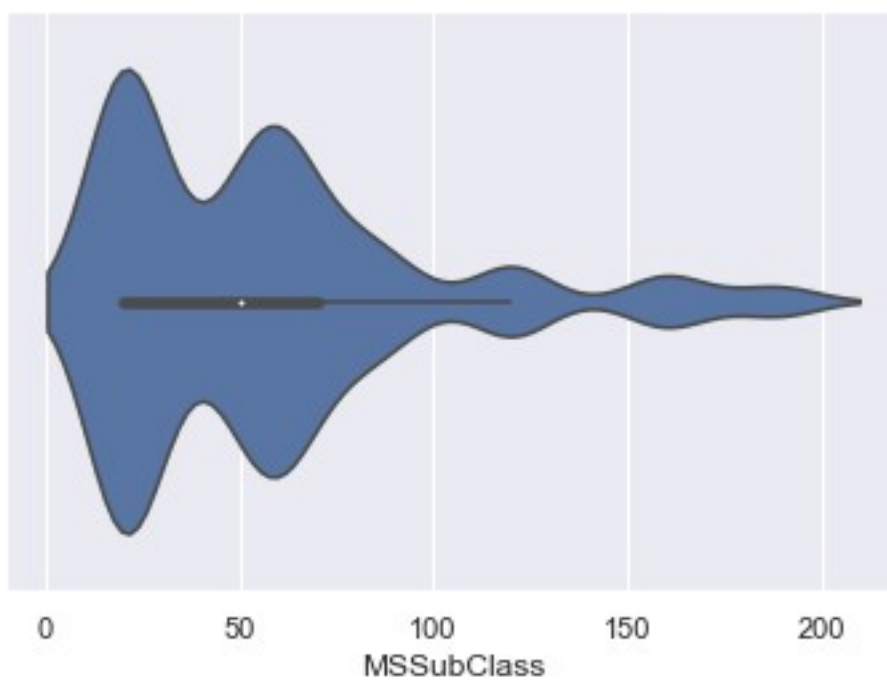

```
Id                0.000165
MSSubClass        1.407011
LotFrontage       1.494021
LotArea          12.573925
OverallQual       0.200786
OverallCond       0.691035
YearBuilt        -0.612295
YearRemodAdd     -0.501838
MasVnrArea        2.696329
BsmtFinSF1        0.764789
BsmtFinSF2        4.251925
BsmtUnfSF         0.920903
TotalBsmtSF       0.511703
1stFlrSF          0.887637
2ndFlrSF          0.812957
LowQualFinSF      9.004955
GrLivArea         1.010992
BsmtFullBath       0.590358
BsmtHalfBath       4.100114
FullBath           0.031271
HalfBath           0.680051
BedroomAbvGr      0.212325
KitchenAbvGr      4.484883
TotRmsAbvGrd      0.660502
Fireplaces        0.632060
GarageYrBlt      -0.693237
GarageCars        -0.342377
GarageArea        0.131748
WoodDeckSF        1.545805
OpenPorchSF       2.339829
EnclosedPorch     3.087164
3SsnPorch        10.297106
ScreenPorch       4.118929
PoolArea          15.948945
MiscVal           24.460085
MoSold            0.215432
YrSold            0.095420
SalePrice         1.881296
TotalSqft         0.804321
TotalBath         0.246687
HouseAge          0.607894
dtype: float64
```


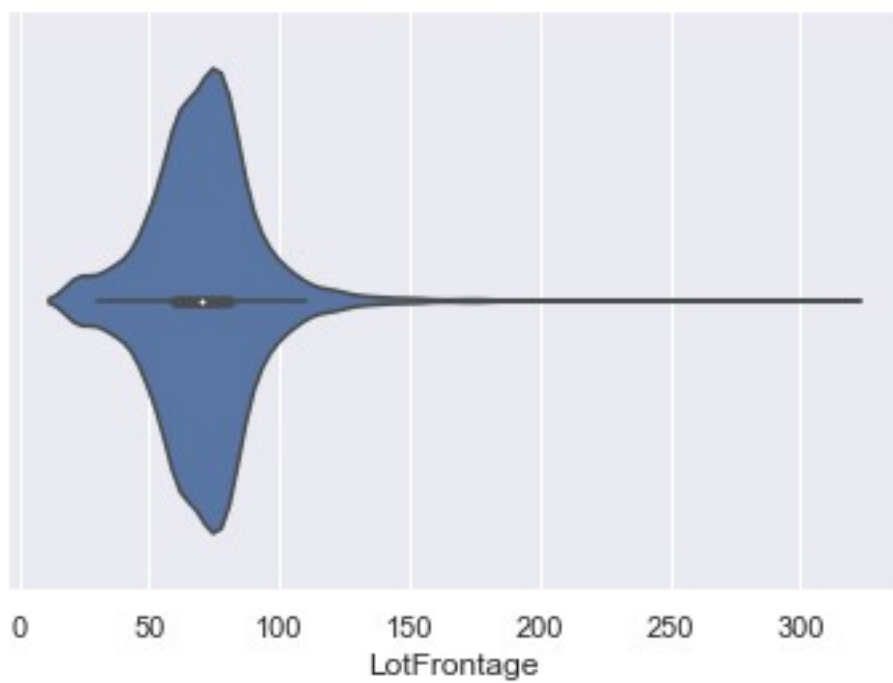
```
cols = dataset_numeric.columns
```



```
for c in cols:
    sns.violinplot(x=dataset_numeric[c])
```


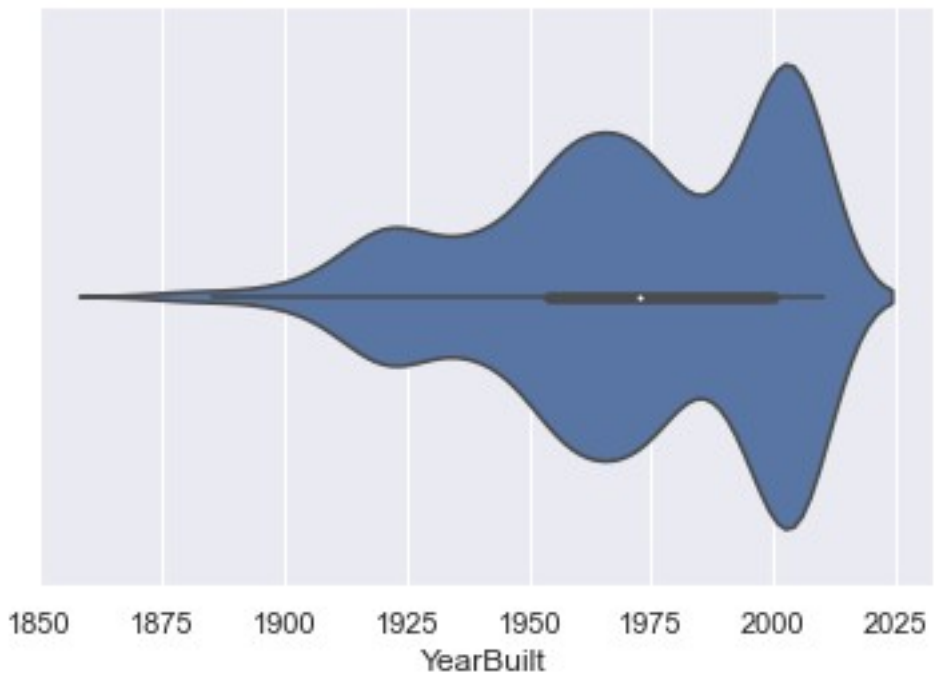



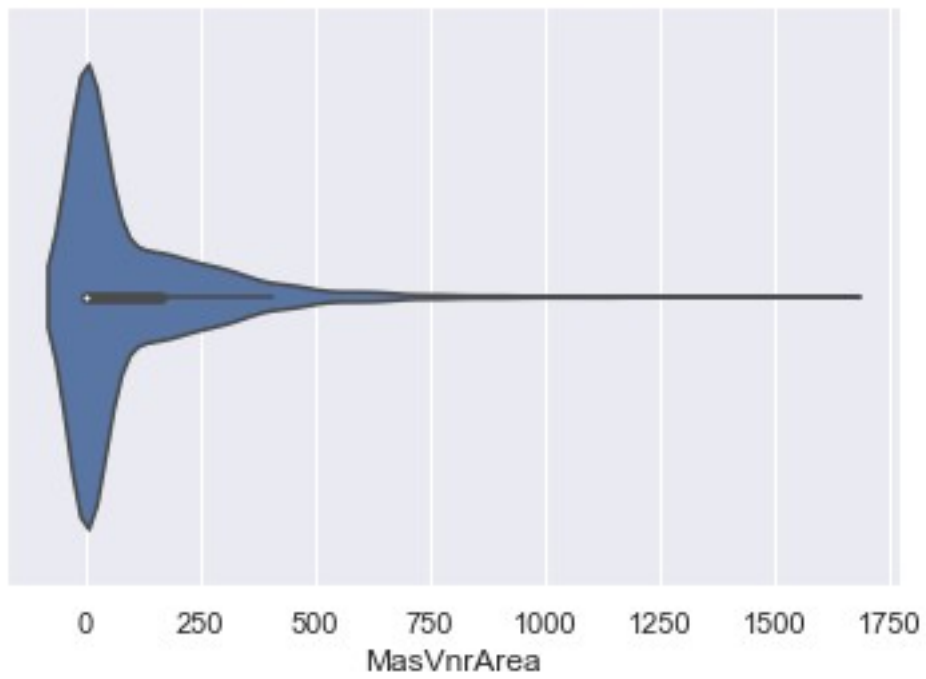
```
plt.xlabel(c)
plt.show()
```









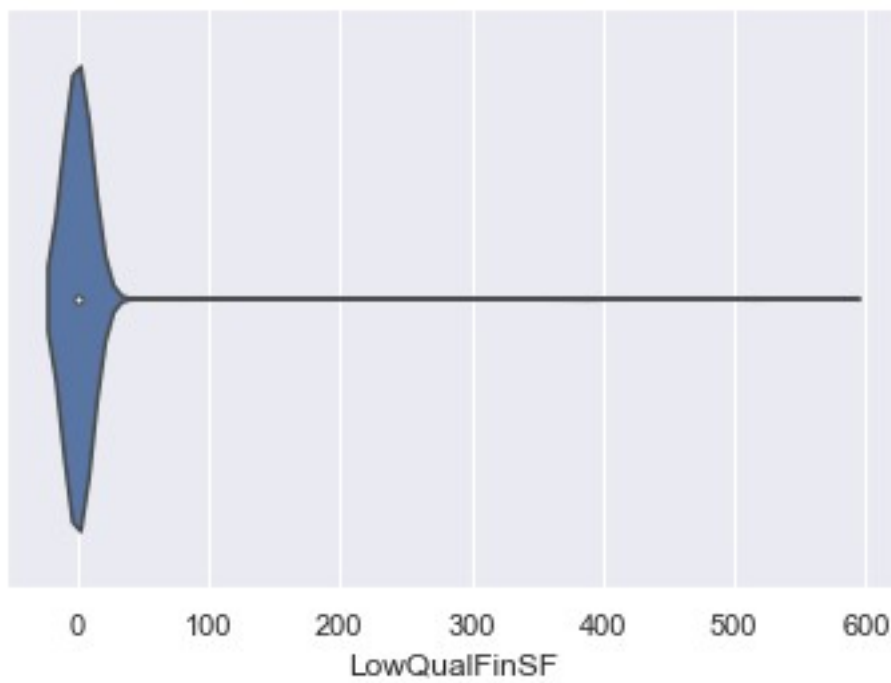
































































The violin plots above reveal several negative (left-sided) skews. Normalization of these distributions can benefit the model.

```
# Skew Correction
#loglp function applies log(1+x) to all elements of the column
skew = df_train.select_dtypes(include=numerics).skew()
```

```
# disabling the pandas warning
pd.options.mode.chained_assignment = None
```

```
skewedfeatures = [s for s in skew if(s > 5.0)]
skewedfeatures
```

```
for skf in skewedfeatures:
    sk = skew[skew == skf].index[0]
    df_train[sk] = np.loglp(df_train[sk])
```

```
# Skew Correction for test set
#loglp function applies log(1+x) to all elements of the column
skew = df_test.select_dtypes(include=numerics).skew()
```

```
# disabling the pandas warning
pd.options.mode.chained_assignment = None
```


```
skewedfeatures = [s for s in skew if(s > 5.0)]
skewedfeatures
```

```
for skf in skewedfeatures:
```

```
sk = skew[skew == skf].index[0]
df_test[sk] = np.log1p(df_test[sk])
```

Prep (Numeric Variables)

```
plt.figure(figsize=(10,8))
sns.heatmap(dataset_numeric.corr(), center = 0)
plt.title("Numeric Correlations")
plt.show()
```



```
# Calculates pearson co-efficient for all combinations
data_corr = dataset_numeric.corr()
```

```
# Set the threshold to select only highly correlated attributes
threshold = 0.5
```

```
# List of pairs along with correlation above threshold
corr_list = []
```

```
size = 36
```

```

#Search for the highly correlated pairs
for i in range(0,size): #for 'size' features
    for j in range(i+1,size): #avoid repetition
        if (data_corr.iloc[i,j] >= threshold and data_corr.iloc[i,j] <
1) or (data_corr.iloc[i,j] < 0 and data_corr.iloc[i,j] <= -threshold):
            corr_list.append([data_corr.iloc[i,j],i,j]) #store
correlation and columns index

```

```

#Sort to show higher ones first
s_corr_list = sorted(corr_list,key=lambda x: -abs(x[0]))

```

```

#Print correlations and column names
for v,i,j in s_corr_list:
    print ("%s and %s = %.2f" % (cols[i],cols[j],v))

```

```

GarageCars and GarageArea = 0.89
YearBuilt and GarageYrBlt = 0.84
GrLivArea and TotRmsAbvGrd = 0.83
TotalBsmtSF and 1stFlrSF = 0.80
2ndFlrSF and GrLivArea = 0.69
BedroomAbvGr and TotRmsAbvGrd = 0.68
BsmtFinSF1 and BsmtFullBath = 0.66
GrLivArea and FullBath = 0.64
GarageYrBlt and GarageCars = 0.62
2ndFlrSF and TotRmsAbvGrd = 0.61
2ndFlrSF and HalfBath = 0.61
YearRemodAdd and GarageYrBlt = 0.60
GarageYrBlt and GarageArea = 0.60
OverallQual and GarageCars = 0.60
YearBuilt and YearRemodAdd = 0.59
OverallQual and GrLivArea = 0.59
OverallQual and YearBuilt = 0.57
OverallQual and GarageArea = 0.56
OverallQual and GarageYrBlt = 0.55
FullBath and TotRmsAbvGrd = 0.55
OverallQual and YearRemodAdd = 0.55
OverallQual and FullBath = 0.55
OverallQual and TotalBsmtSF = 0.54
GrLivArea and BedroomAbvGr = 0.54
YearBuilt and GarageCars = 0.54
1stFlrSF and GrLivArea = 0.53
BsmtFinSF1 and BsmtUnfSF = -0.52
2ndFlrSF and BedroomAbvGr = 0.50

```

Prep (Categorical Variables)

```

dum_vars = ['Neighborhood', 'MSZoning', 'MSSubClass', 'Street',
            'LotShape', 'LotConfig', 'Utilities', 'LandSlope',
            'BldgType', 'HouseStyle', 'RoofStyle', 'Foundation',

```

```
    'Heating', 'CentralAir', 'PavedDrive', 'MoSold',  
    'YrSold', 'SaleType', 'SaleCondition']
```

```
for col in dum_vars:  
    print(col)  
    train = sorted(df_train[col].unique().tolist())  
    test = sorted(df_test[col].unique().tolist())  
    total = set(train + test)  
    df_train[col] = pd.Categorical(df_train[col], categories=total)  
    df_test[col] = pd.Categorical(df_test[col], categories=total)
```

Neighborhood
MSZoning
MSSubClass
Street
LotShape
LotConfig
Utilities
LandSlope
BldgType
HouseStyle
RoofStyle
Foundation
Heating
CentralAir
PavedDrive
MoSold
YrSold
SaleType
SaleCondition

Models

Lasso Regression

"Least Absolute Shrinkage and Selection Operator" Using Lasso linear regression we will "shrink" values towards the mean to simplify the model. This form of regression will hopefully help reduce the impact of noise on our model.

Feature(s) to look at

```
f1 = ['MSSubClass', 'LotFrontage', 'LotArea',  
      'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',  
      'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',  
      'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath',  
      'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr',  
      'TotRmsAbvGrd', 'Fireplaces', 'GarageArea', 'WoodDeckSF',  
      'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch',  
      'PoolArea', 'MiscVal', 'MoSold', 'YrSold',  
      'TotalSqft', 'TotalBath', 'HouseAge']
```



```

# Run a Linear Regression using the feature(s)
x1 = df_train[f1]
y = df_train['SalePrice']

# Split the data
x_train, x_test, y_train, y_test = train_test_split(x1, y)
x_train.shape, x_test.shape, y_train.shape, y_test.shape

((1093, 34), (365, 34), (1093,), (365,))

# Set up model
lasso = LassoCV(n_alphas=200, alphas=np.logspace(0, 4, 100), max_iter
= 11000)
kf = KFold(n_splits=7, shuffle=True)

# Standardize the data
ss = StandardScaler()
ss_train = ss.fit_transform(x_train)
ss_test = ss.transform(x_test)

# cross validate
scores = cross_val_score(lasso, ss_train, y_train, cv=kf)
print(scores)
print(f'Mean Score: {scores.mean()}; SD: {scores.std()}')

lasso.fit(ss_train, y_train)
print(f'TRAIN Score: {lasso.score(ss_train, y_train)}')
print(f'TEST Score: {lasso.score(ss_test, y_test)}')

pred = lasso.predict(ss_test)
b, m = np.polynomial.polynomial.polyfit(y_test, pred, 1)

[0.86678082 0.87698432 0.83635709 0.82759973 0.8937877 0.88209416
0.79524837]
Mean Score: 0.8541217392587718; SD: 0.03278472706223093
TRAIN Score: 0.864825923973606
TEST Score: 0.8617164697332577

```

My initial attempt at running this model (with max_iter = 1000) yeilded a ConvergenceWarning. To avoid this warning, I decided to drastically increase the iterization. However, this is computationally inefficient and could negatively impact the model. The initial warnings could be a sign that the model is not fitting to the data correctly. That being said, we are left with a bit of confusion considering the mean score seems to reflect a well optimized model.

```

# Visualize the model results
sns.scatterplot(x=y_test, y=pred, alpha=0.4)
sns.regplot(x=y_test, y=pred, truncate=True, scatter_kws={'s': 20,

```

```

'alpha':0.3},
    line_kws={'color':'red', 'linewidth': 2})
sns.lineplot(x=np.unique(y_test), y=np.unique(np.poly1d(b + m *
np.unique(y_test))), linewidth=0.5, color='r')

plt.xlabel("Actual Sale Prices")
plt.ylabel("Predicted Sale Prices")
plt.title("Actual Prices vs Predicted prices [Test Set]")

plt.show()

```



Prediction Score

```

holdout_df = df_test[f1]

# Standardize the numeric columns
ss = StandardScaler()
ss_holdout = ss.fit_transform(holdout_df)

# predict SalePrice
predict = lasso.predict(ss_holdout)
submit = pd.DataFrame({'Id': df_test['Id'], 'SalePrice': predict})
submit

#export to csv
submit.to_csv('lasso_submission.csv',index=False)

```

```
from IPython import display
display.Image("lasso_score.png")
```

YOUR RECENT SUBMISSION



lasso_submission.csv

Submitted by SeafoodTakeout · Submitted just now

Score: 0.21471

Not a great score, but a worthy first attempt.

Ridge Regression

Ridge regression aims to shrink the dimensions of the data by equal proportions.

```
# Feature(s) to look at
f2 = ['MSSubClass', 'LotFrontage', 'LotArea',
      'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
      'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',
      'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath',
      'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr',
      'TotRmsAbvGrd', 'Fireplaces', 'GarageArea', 'WoodDeckSF',
      'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch',
      'PoolArea', 'MiscVal', 'MoSold', 'YrSold',
      'TotalSqft', 'TotalBath', 'HouseAge']

# Run a Linear Regression using the feature(s)
x2 = df_train[f2]
y = df_train['SalePrice']

# Split the data
x2_train, x2_test, y2_train, y2_test = train_test_split(x2, y)
x2_train.shape, x2_test.shape, y2_train.shape, y2_test.shape

((1093, 34), (365, 34), (1093,), (365,))

# Set up model
ridge_reg = Ridge(alpha=1, solver="cholesky")
kf = KFold(n_splits=7, shuffle=True)

# Standardize the data
ss = StandardScaler()
ss_train2 = ss.fit_transform(x2_train)
ss_test2 = ss.transform(x2_test)

# cross validate
scores = cross_val_score(ridge_reg, ss_train2, y2_train, cv=kf)
print(scores)
print(f'Mean Score: {scores.mean()}; SD: {scores.std()}')

ridge_reg.fit(ss_train2, y2_train)
```

```

print(f'TRAIN Score: {ridge_reg.score(ss_train2, y2_train)}')
print(f'TEST Score: {ridge_reg.score(ss_test2, y2_test)}')

pred2 = ridge_reg.predict(ss_test2)
b, m = np.polynomial.polynomial.polyfit(y2_test, pred2, 1)

[0.77746177 0.85340441 0.88977111 0.8172674 0.85225053 0.87785206
 0.82413569]
Mean Score: 0.8417347086521818; SD: 0.03565330722620647
TRAIN Score: 0.859717191911943
TEST Score: 0.885335525073244

```

Fortunately, running the ridge regression model did not prompt any warnings. It appears to have fit to the data fairly well, however, the mean score is slightly lower than that of the 11000 iteration Lasso model.


```

# Visualize the model results
sns.scatterplot(x=y2_test, y=pred2, alpha=0.4)
sns.regplot(x=y2_test, y=pred2, truncate=True, scatter_kws={'s': 20,
'alpha':0.3},
            line_kws={'color':'red', 'linewidth': 2})
sns.lineplot(x=np.unique(y2_test), y=np.unique(np.poly1d(b + m *
np.unique(y2_test))), linewidth=0.5, color='r')

plt.xlabel("Actual Sale Prices")
plt.ylabel("Predicted Sale Prices")
plt.title("Actual Prices vs Predicted prices [Test Set]")

plt.show()

```



Prediction Score

```
holdout_df2 = df_test[f2]
```

```
# Standardize the numeric columns
```

```
ss = StandardScaler()
```

```
ss_holdout2 = ss.fit_transform(holdout_df2)
```

```
# predict SalePrice
```

```
predict2 = ridge_reg.predict(ss_holdout2)
```

```
submit2 = pd.DataFrame({'Id': df_test['Id'], 'SalePrice': predict2})
```

```
submit2
```

```
#export to csv
```

```
submit2.to_csv('ridge_submission.csv', index=False)
```

```
display.Image("ridge_submission.png")
```

YOUR RECENT SUBMISSION



ridge_submission.csv

Submitted by SeafoodTakeout · Submitted just now

Score: 0.34951

This is not a great score. It is significantly worse than the Lasso model and leaves me with a lot of questions and concerns about the initial quality of my prep on the dataset.

Elastic Net

Elastic Net regression, in a way, fills the middle ground between lasso and ridge regression. It uses both L1 and L2 regularization techniques to try and capture the benefits while offsetting each others potential pitfalls.

```
# Feature(s) to look at
f3 = ['MSSubClass', 'LotFrontage', 'LotArea',
      'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
      'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',
      'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath',
      'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr',
      'TotRmsAbvGrd', 'Fireplaces', 'GarageArea', 'WoodDeckSF',
      'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch',
      'PoolArea', 'MiscVal', 'MoSold', 'YrSold',
      'TotalSqft', 'TotalBath', 'HouseAge']

# Run a Linear Regression using the feature(s)
x3 = df_train[f3]
y = df_train['SalePrice']

# Split the data
x3_train, x3_test, y3_train, y3_test = train_test_split(x3, y)
x3_train.shape, x3_test.shape, y3_train.shape, y3_test.shape

((1093, 34), (365, 34), (1093,), (365,))

# Set up model
e_net = ElasticNetCV(n_alphas=200, alphas=np.logspace(0, 4, 100),
max_iter = 1000)
kf = KFold(n_splits=7, shuffle=True)

# Standardize the data
ss = StandardScaler()
ss_train3 = ss.fit_transform(x3_train)
ss_test3 = ss.transform(x3_test)

# cross validate
scores = cross_val_score(e_net, ss_train3, y3_train, cv=kf)
print(scores)
print(f'Mean Score: {scores.mean()}; SD: {scores.std()}')

e_net.fit(ss_train3, y3_train)
print(f'TRAIN Score: {e_net.score(ss_train3, y3_train)}')
print(f'TEST Score: {e_net.score(ss_test3, y3_test)}')

pred3 = e_net.predict(ss_test3)
b, m = np.polynomial.polynomial.polyfit(y3_test, pred3, 1)
```

```
[0.86790248 0.84501076 0.8735932 0.80718903 0.81116816 0.81723714  
0.82687952]
```

```
Mean Score: 0.8355686140825493; SD: 0.025050885363058557
```

```
TRAIN Score: 0.8485104942436186
```

```
TEST Score: 0.8478841103210351
```

```
# Visualize the model results
```

```
sns.scatterplot(x=y3_test, y=pred3, alpha=0.4)  
sns.regplot(x=y3_test, y=pred3, truncate=True, scatter_kws={'s': 20,  
'alpha':0.3},
```


```
line_kws={'color':'red', 'linewidth': 2})  
sns.lineplot(x=np.unique(y3_test), y=np.unique(np.poly1d(b + m *  
np.unique(y3_test))), linewidth=0.5, color='r')
```

```
plt.xlabel("Actual Sale Prices")
```

```
plt.ylabel("Predicted Sale Prices")
```

```
plt.title("Actual Prices vs Predicted prices [Test Set]")
```

```
plt.show()
```



The Elastic Net model ran without complications. However, based on the mean score and the standard deviation it appears to have fit to the data worse than the previous two models. That being said, it still fits appropriately and overfitting won't seem to be a problem.

```
Prediction Score
```

```
holdout_df3 = df_test[f3]
```

```

# Standardize the numeric columns
ss = StandardScaler()
ss_holdout3 = ss.fit_transform(holdout_df3)

# predict SalePrice
predict3 = e_net.predict(ss_holdout3)
submit3 = pd.DataFrame({'Id': df_test['Id'], 'SalePrice': predict3})
submit3

#export to csv
submit3.to_csv('elasticnet_submission.csv', index=False)

display.Image("elasticnet_score.png")

```

YOUR RECENT SUBMISSION



elasticnet_submission.csv
Submitted by SeafoodTakeout · Submitted just now

Score: 0.16614

The Elastic Net submission score was by far the best of the three. Despite its initial fit, the model seemed to handle noise well. The combination of L1 and L2 regularization likely played a synergistic role in the performance of this model.

Conclusion

The three models ran fairly well for first attempts. The Kaggle prediction scores fell within a range of effectiveness. It's likely that their performance could have been optimized by better data prep. However, given the circumstances, these models adequately illustrated many of the performance differences between each regression method. Although, it would be an interesting personal project to see exactly how significantly better prep would optimize these regression methods.